

Project Checkpoint: CSE 6730 Project 1

Chris Dunlap, Allen Koh, Matt May

Spring 2016

Contents

1 Problem Statement	2
2 Related Work	4
3 Conceptual Model	6
3.1 Inputs	6
3.2 Output	8
3.3 Content	9
3.3.1 Approach	9
3.3.2 Parameters	12
3.3.3 Random Numbers and Verification	14
3.4 Assumptions and Simplifications	14
4 Description of Simulation Software	16
4.1 Architecture	16
4.2 Interfaces	16
4.2.1 User Interface	16
4.2.2 Pedestrian Class Interfaces	20
4.2.3 Node Class Interfaces	21
4.2.4 Intersection Class Interfaces	23
5 Results	24
5.1 Verification Methodology	27
5.2 Validation Methodology	27
6 Conclusion	28
7 Final Notes	28

1 Problem Statement

The efficient evacuation of physical structures is an important problem with a long history of multi-disciplinary study [19]. For this study, we focus specifically on the problem of efficiently evacuating the area around Bobby Dodd Stadium in Atlanta, GA. Bobby Dodd Stadium, the home of the NCAA Division-I Georgia Tech Yellow Jackets football team, has a seating capacity of 55,000 individuals.



Figure 1: Bobby Stadium Stadium seating chart.

Our primary objective in this study is to minimize the pedestrian evacuation time of the stadium and its surrounding area as attendees leave the area following a home football game. We do this by taking a parametric

approach, in which we explore the effects of road closures and the “takeover” of certain intersections by law enforcement in promoting an optimal result.

For purposes of this study, we define the system under investigation (SUI) as a rectangular polygon (Fig. 2) surrounding the stadium. To be clear, the SUI is defined as the area inside the polygon but outside the stadium. As pedestrians exit the stadium, they enter the SUI, and as they cross the boundary of the polygon, they exit it.

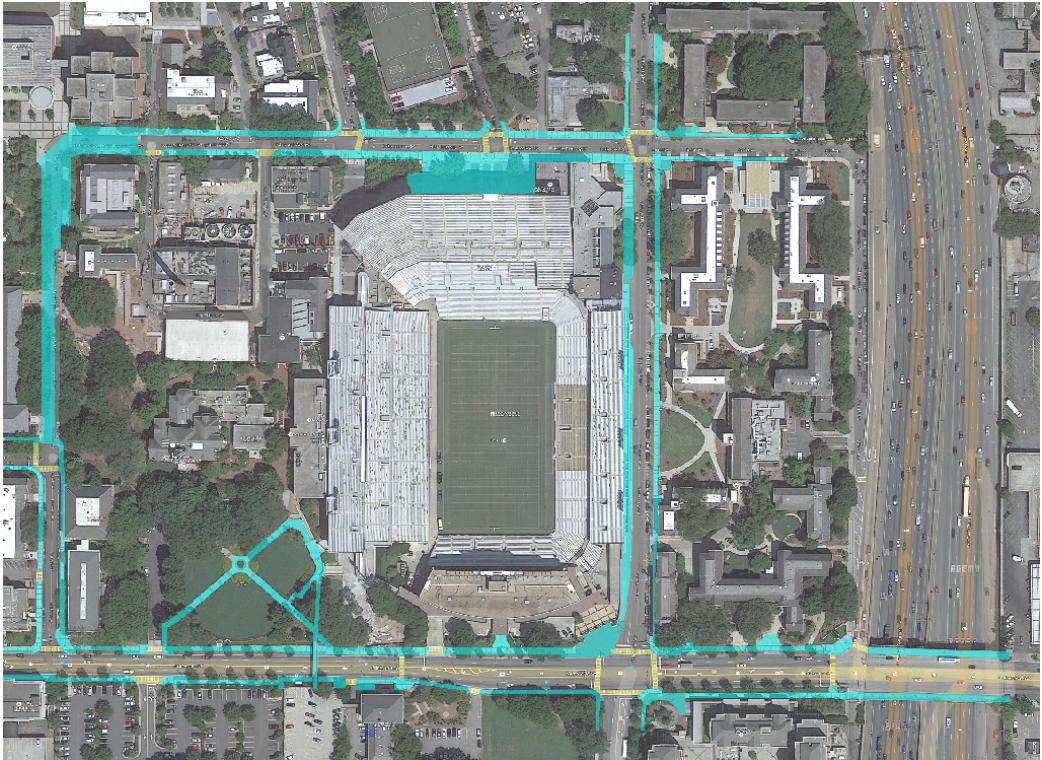


Figure 2: The system under investigation.

The stadium has 11 “exit zones” modeled that serve as ingress (pre-game) and egress (post-game) sites. These serve as the entry points to our simulation. As pedestrians enter the SUI, they are assigned a destination and then proceed to their destination by way of pedestrian walking paths (i.e., sidewalks). These are denoted as the cyan-colored portions of the SUI. Pedestrian crosswalks are also modeled; these are the yellow-colored portions of the SUI.

2 Related Work

A variety of techniques have been used to model pedestrian movement. Cellular automata, lattice gas, social force, fluid-dynamic, agent-based, game-theoretic, and animal experimentation-based approaches have been used [19]. As noted by Zheng, Zhong, and Liu [19], models often encounter or attempt to model several common phenomena: clogging, side-stepping, lane formation, and herding behavior, among them.

These phenomena manifest themselves in different ways and at different magnitudes depending on the system being modeled. In one of the early explorations of 2D cellular automata for simulating traffic flow in the related domain of vehicular traffic simulation, Biham and Middleton [1] note the presence of a sharp *jamming transition* in which all cars in the simulation transition from moving at maximal speed to being stuck. Similar effects were noted in simulations of bi-directional pedestrian movement, with Weifeng, Lizhong, and Weicheng [17] finding that as total pedestrian density increases, a critical value is reached at which the system transitions into a jammed state where only a few pedestrians are able to proceed.

In a slightly different take on the problem, Okazaki and Matsushita [12] modeled pedestrian evacuation using Coulomb's Law, with actors magnetically moving toward their goals and away from obstacles that would lead to collisions. In another application of equations of natural phenomena, Helbing [7] derived fluid dynamic equations for explaining the movement of pedestrian crowds, observing phenomena such as the development of lanes, jamming, and crossing. Helbing [9] further notes an explicit "faster-is-slower" effect, in which attempting to move faster results in a lower average speed of leaving at high pedestrian density in evacuation scenarios.

In a more recent study, Helbing et al. [8] use a "social force" model, which states that pedestrians operate in some sense automatically when reacting to obstacles and other pedestrians, applying strategies that have been learned to be most effective over time. Several suggestions are made to alleviate the three most common problems in pedestrian crowds: counterflows, bottlenecks, and intersecting flows. The presence of strategically placed obstacles are found to actually reduce these negative phenomena, leading to improved flow.

Building on previous cellular automata-based models, Burstedde et al. [3] introduce a *floor field*, a secondary grid of cells which underlies the main grid and acts as a substitute for pedestrian intelligence. These fields can be either

static or dynamic, and are capable of promoting the avoidance of jams, as well as simulating attractive effects, in which pedestrians are more likely to follow in the paths of pedestrians ahead of them.

Taking a more algorithmic approach, Fang et al. [6] have designed a modified ant colony optimization (ACO) algorithm for optimizing pedestrian evacuation, seeking to minimize evacuation time, evacuation distance, and congestion. Kemloh Wagoum, Seyfried, and Holl [10] utilize an observation principle approach in modeling pedestrian evacuation, in which pedestrians first observe their environment and then make a final decision on strategy based on obtained data.

There are several takeaways from the literature described above that have potential applicability to the model at hand. First, on the topic of modeling pedestrian walkways, several papers focused on either a single passage-way, enclosed areas with obstacles (walls, pillars, etc.), or evacuation from a room with a single doorway. However, modeling walkways as an undirected graph [6] seems feasible given the problem at hand, where there are multiple pathways one can take to get to a single destination. In papers where cellular automata was utilized, the median size of each cell on the walkway was 0.4m^2 [2, 3, 17].

Second, on the topic of pedestrian modeling, multiple walking speed techniques were used, ranging from a constant speed of 1 m/s [17] or 1.3 m/s [3], to using a distribution resembling a step function [2] or a Normal distribution [11]. In either distribution, the median walking speed was also approximately 1.3 m/s. For reaction to stimuli, 0.3 seconds was utilized [3], which was one time step in that particular simulation. Incorporation of opportunistic side-stepping was done [2], in addition to using least-cost algorithms for establishing the path a pedestrian would like to take [6]. In the context of cellular automata, analyzing a given pedestrian's cell's neighbors in either a 4-connected [17] or 8-connected fashion [3] to decide in which direction the pedestrian would like to walk for the next time step was utilized. This, coupled with a dynamic floor field [3] to factor in long-distance forces, could be utilized. In cases where there are two pedestrians targeting each other's current positions for a next move, which may happen when they are walking in opposing directions in the same lane, the concept of *place exchange* has merit when two pedestrians walking in opposing directions [2].

Finally, simulation execution practices were also gleaned. Care must be taken to do parallel-based updating to avoid sequential updating from interfering with underlying model execution [2]. The practice of using a deter-

ministic model running with randomized initial conditions has precedent [1]. Finally, to achieve statistical significance, 20 runs per configuration has been used [2].

3 Conceptual Model

For our simulation, we utilize a 2-dimensional cellular automata (CA)-based approach. We construct a discrete time-stepped model, in which the simulation clock advances by a fixed interval every time step. In our simulation, 1 time step is modeled as 1 second. We utilize a stochastic (probabilistic) approach, in which the output of each individual simulation run is a random variable. Thus, multiple simulation runs will be performed and statistical analysis will be used to provide confidence levels in our results.

3.1 Inputs

The primary input to our simulation model is a stream of football game attendees (i.e., pedestrians) exiting Bobby Dodd Stadium following a home football game. For purposes of modeling this input stream, we assume pedestrian interarrival time at the stadium gate boundary is a homogeneous stochastic process that follows a Poisson distribution, as has been used in the literature [16]. Empirically, we determined the mean number of pedestrians per second arriving at the exit boundary from a 5-meter-wide exit to a stadium to be 4.05. This was determined by observing the number of pedestrians exiting a stadium on video [15]. This value was then used as the basis for the lambda-parameter used for sampling of the Poisson distribution. This provides a stream of exiting pedestrians in which the number of pedestrians exiting at any given second is independent of the pedestrians exiting at the previous second. While this assumption is reasonable for an initial simulation model, it should be noted that in reality there is likely some variation over time. We use a simple queue structure to model the entry of pedestrians into our simulation. As pedestrians arrive at the gate boundary and exit the stadium, they enter the SUI.

The random number generator used to sample the Poisson distribution underwent a chi-squared goodness-of-fit test to evaluate the hypothesis that the numbers obtained were indeed following a Poisson distribution. A 7-bin histogram of 1771 observed values from the generator, along with the

corresponding theoretical histogram, is shown in Fig. 3. The chi-squared statistic was 1.2640; when compared with the critical value of 11.0705 (from chi-squared distribution with 5 degrees of freedom and 0.05 level of significance), the hypothesis stating that the distribution is Chi-Squared is not rejected.




Figure 3: Poisson PDF Chi-Squared Test ($\lambda = 82.4280$).

In addition, we select each pedestrian's target destination and walking speed at random. For selecting a target destination, we assume that pedestrian destinations are uniformly distributed and sample from a pool of exit nodes from the SUI that correspond to parking lots, local housing, and public transportation. For determining a walking speed for each pedestrian, we sample from a distribution previously referenced in the literature by Blue and Adler [2] (5% fast walkers (4 cells/time step), 90% standard (3 cells/time step), 5% slow (2 cells/time step)).

Pedestrians in our model can be considered a consumer entity class *Pedestrian* with attributes which will be set and updated during the course of our simulation. Table 1 provides an overview of the attributes of the *Pedestrian*

class.

Attribute	Description
Current	<i>Node</i> object corresponding to the current location of the pedestrian.
Destination	<i>Node</i> object (see Table 2) corresponding to the final destination of the pedestrian.
TargetNext	The desired next node to move to in the pedestrian’s shortest path to his destination, as computed by Dijkstra’s algorithm.
Speed	Walking speed, formulated in grid cells traversed per time step.
EgressComplete	A boolean value. This value is false when the pedestrian has not yet egressed, and true when egress is complete.
ShortestPath	List of node ids, corresponding to the shortest path to the pedestrian’s destination.

Table 1: Attributes of the *Pedestrian* entity class.

New instances of the *Pedestrian* class are created and initialized as pedestrians exit the stadium.

3.2 Output

For each simulation run, when all pedestrians are evacuated from the SUI we determine the simulation to be complete. As this is a stochastic simulation, the output results must be interpreted as samples of a random process.

At the conclusion of a series of simulation runs for a given set of parameters, we compute an average of the total egress duration across all runs from that series. More formally, the average egress time E_{avg} for n simulation runs can be represented as:

$$E_{avg} = \frac{1}{n} \sum_{i=1}^n d_i \quad (1)$$

where d_i is the total egress duration of the i th simulation run. This output can be considered a derived scalar output variable (DSOV). We then define the optimal parameter strategy to be the set of parameters P that give the minimum value of E_{avg} . Confidence intervals of 90% are also computed to provide context on the probable range of E_{avg} across all parameters in P .

3.3 Content

3.3.1 Approach

Pedestrians enter the SUI from one of 11 potential entrance zones, shown in Fig. 4. To build the simulation space, we overlay a 2-dimensional cellular automata grid on a map of the SUI. We do this by modeling the SUI as an undirected weighted graph. In our case, the edge weights correspond to the spatial distance between two nodes [18]. Note that in this section, the terms “cell” and “node” are used interchangeably.

As pedestrians exit the stadium (and thus enter the SUI), they are probabilistically assigned one of the graph’s destination nodes (exit nodes in Fig. 4) according to a uniform distribution. When a pedestrian reaches a destination node, that constitutes exiting the SUI. Destination nodes were determined based on their spatial meaning when overlaid on a map of the SUI, and constitute one of three possible classes: on-campus housing (dormitories), parking lots, or the North Avenue Metropolitan Atlanta Rapid Transit Authority (MARTA) station.

Each node in the graph is a member of a *Node* class, which has attributes as shown in Table 2. As part of the simulation initialization, the node and edge information is populated. We define a few key types of nodes, such as entrance, exit (or destination), and sidewalk nodes. For each entrance node N_e in the graph, we precompute the least-cost path to every possible destination node N_d . In general, we can compute the shortest path between any given node N_i and destination node N_d using Dijkstra’s well-known algorithm [4]. In our simulation, we use a publicly available Python implementation of Dijkstra’s algorithm [5]. As pedestrians are initialized and added to the simulation, they are randomly assigned an entrance node according to a uniform distribution. The pedestrian then “looks up” the appropriate path to the assigned destination. This path is then stored as a list structure on the pedestrian called *ShortestPath*.

At each time step, the simulation pops an element (which corresponds

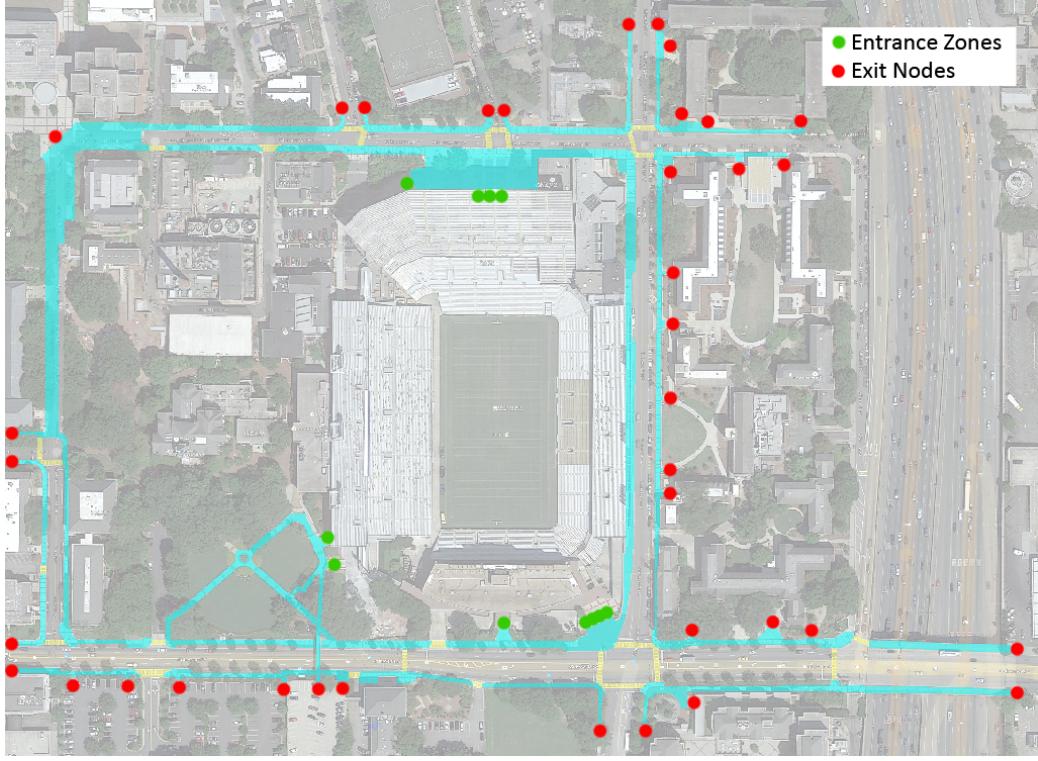


Figure 4: SUI Entrances and Exits.

to a node) from each active pedestrian’s *ShortestPath* attribute, and then attempts to move the pedestrian to the given node. This repeated process drives the pedestrian’s instinct to move toward their destination, based strictly on the shortest distance walk from their current position to their destination node.

It will likely occur that more than one pedestrian Ped_A will decide to pursue a node N_B that is already occupied by a pedestrian Ped_B . In such cases, a check is first performed to see if Ped_B is attempting to move to the location occupied by Ped_A ; if so, the *place exchange* logic [2] is activated, and the pedestrians trade places. However, supposing the *place exchange* logic is not triggered, then an empty neighbor N_E of Ped_A is chosen as the next node to move to instead. Ped_A moves to N_E and the *ShortestPath* list of Ped_A is expanded to include the steps that must be taken to get from N_E to N_B , such that the path from N_B to the final destination remains intact. If there are no empty neighbors, then Ped_A does not move at all.

Attribute	Description
NodeID	Unique identifier for the node.
XCoordinate	X-coordinate of the node on the 2-dimensional grid.
YCoordinate	Y-coordinate of the node on the 2-dimensional grid. (Note: for purposes of this simulation, the y-axis is inverted.)
PixelX	X-coordinate of the node in 2-dimensional pixel space, for plotting on top of a map of the SUI.
PixelY	Y-coordinate of the node in 2-dimensional pixel space, for plotting on top of a map of the SUI.
Paths	A dictionary (hash map) relating each possible destination to the shortest path as determined by Dijkstra's algorithm.
Neighbors	Array of <i>Node</i> objects corresponding to nodes that are one edge away from the node.
Available	Boolean value indicating whether the node is available for a pedestrian (true) or not available (false).
NodeType	Enumeration of the type of node, corresponding to its spatial significance. The NodeType can be one of four possibilities: sidewalk, road, SUI entrance, or SUI exit.
CurrentPed	Pedestrian currently occupying the node, if any.

Table 2: Attributes of the *Node* class.

3.3.2 Parameters

To optimize the egress of pedestrians from the SUI, there are several parameters that can be altered. The first parameter is opening certain streets or intersections meant for vehicular traffic and making them for use by pedestrians only (i.e., roads become sidewalks). The second is closing certain streets or intersections and making them for use by vehicles only (i.e. intersections become roads). The optimal time for all pedestrians to evacuate the SUI could be a combination of the two parameters. For example, opening an intersection to pedestrians only and closing a nearby small intersection may to pedestrian traffic (i.e., opening it to vehicular traffic only) may be a combination that produces better results than only directing pedestrians to a regular crosswalk. We explore these types of permutations in our results.

To simulate these parameters, in addition to a *Node* class, there is also an *Intersection* class, which has attributes as shown in Table 3.

Attribute	Description
IntId	Integer value. A unique identifier for the intersection.
Nodes	Array of <i>Node</i> objects corresponding to nodes that are part of this intersection.
IsOpen	Boolean value indicating whether the intersection is open for pedestrian access (true) or not open (false).

Table 3: Attributes of the *Intersection* class.

Intersections are modeled by mapping a given intersection ID (*IntId*) to the nodes that comprise it. The intersection is therefore made up of a group of nodes. To simulate the crosswalk opening and closing, the nodes will be either forced closed or available for pedestrians to occupy them. If the road is closed and vehicles are passing, then no nodes in the intersection will be available for pedestrians.

One parameter in the simulation that can be modified is determining which intersections will always be open to pedestrian use (i.e., closed to vehicles). To model this case, all nodes in that intersection will behave in the same way as the nodes that make up the sidewalk. There will be no need

to create an intersection object for that intersection. On the other hand, if there is an intersection that is closed off for pedestrians and only available for vehicular traffic, then the initial creation of nodes will not include that intersection - there will be no nodes at that intersection as no pedestrian can ever occupy those nodes. For normal intersections, an intersection object is created and a time is assigned to the intersection which defines the number of simulation timesteps before changing the state of the intersection (open or closed). Fig. 5 displays the intersections and their corresponding IDs. Table 4 defines the open/close timesteps for each intersection. These times were selected based on the size of the intersections and the importance of the roads associated with those intersections.

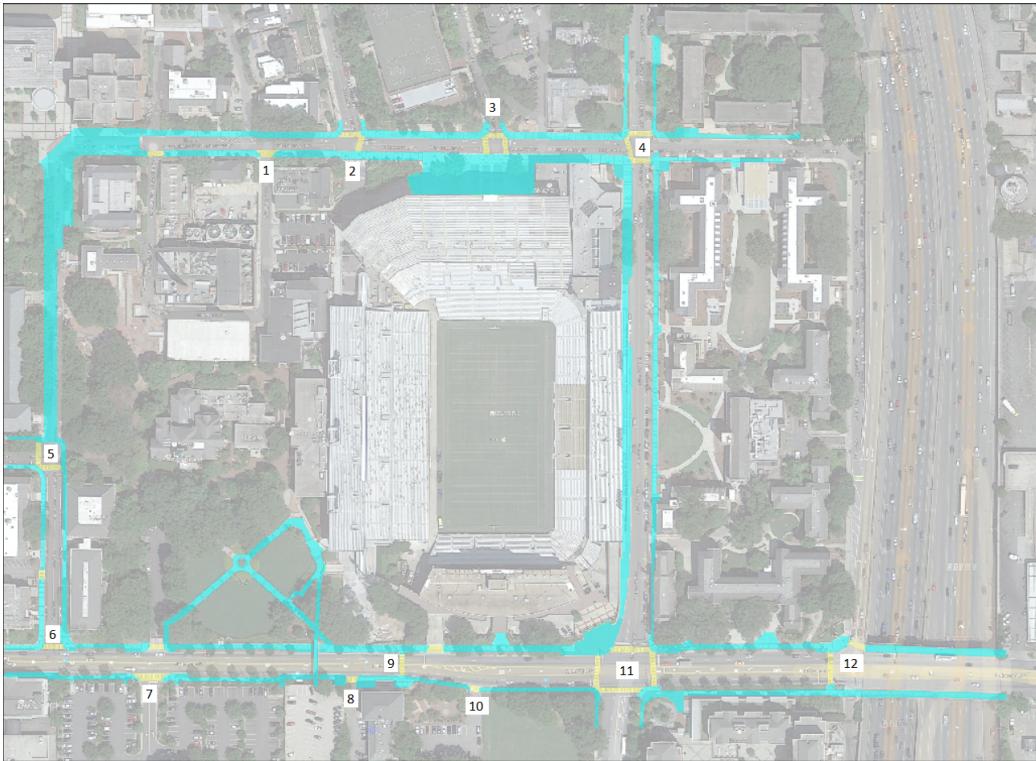


Figure 5: SUI Intersection IDs.

If there are pedestrians that need to cross a normal intersection, but the intersection is closed, then the pedestrians will not be able to advance to their next node. This will cause a list of pedestrians to stay in their current

Intersection ID	1	2	3	4	5	6	7	8	9	10	11	12
Timesteps	15	20	25	30	25	20	15	15	20	15	40	25

Table 4: Intersection Open/Close Timesteps

node until the crosswalk opens up (i.e., implicitly form a queue). Once that occurs, then the pedestrians can continue to their next nodes.

3.3.3 Random Numbers and Verification

To model the egress of the stadium, we take a stochastic (probabilistic) approach. In the context of a computer simulation, this necessitates the use of a pseudorandom number generator.

For our simulation, we use the Lehmer random number generator [14], also known as the Park-Miller random number generator [13]. For verification of the generator, we performed 100,000 iterations of a chi-square test procedure in which 1,000 samples were drawn from the random number generator. For each iteration, we placed each sample into one of 100 bins, forming a histogram. We then computed the chi-square statistic, namely:

$$\chi^2 = \sum_{i=1}^n \frac{(O_i - E_i)^2}{E_i} \quad (2)$$

where χ^2 is the chi-squared statistic, which asymptotically approaches a χ^2 distribution, O_i is the count of observations of type i , and E_i is the expected count of observations of type i . We determined the chi-square critical value for 99 degrees of freedom and a p -value of 0.05 to be 123.225. Experimentally, our generator failed to pass the chi-square test only 4.95% of the time over 100,000 iterations, allowing us to not reject the null hypothesis that our random number generator produces uniformly distributed random numbers with 95% certainty.

3.4 Assumptions and Simplifications

We make a number of assumptions and simplifications in constructing our simulation model.

Firstly, our simulation is simplified by focusing only on *pedestrian* traffic, avoiding the inclusion of vehicular traffic which could be significant in a football game evacuation scenario.

We also assume the stadium to be at or near full seating capacity of 55,000 individuals. While the mean empirical value of football game attendance at Bobby Dodd Stadium is somewhat lower than 55,000, the choice of 55,000 allows us to “stress test” our model by simulating the upper bound.

We assume that pedestrians stay on sidewalks and obey traffic lights, as the inclusion of other types of pedestrian behaviors would likely add complexity to the model without yielding much additional value. We also assume that main roads such as North Avenue remain open to vehicular traffic, avoiding a more unrealistic simulation that would be possible if all roads could be closed.

In the realm of our “updating rule” which specifies the next cell for a given pedestrian to travel to in the simulation, we assume that pedestrians are following a shortest-path approach. While logical, this may be an oversimplification, as pedestrians (particularly from the home team) may follow more indirect paths to their destinations following a victory.

We make the assumption that all pedestrians will be traveling to one of three possible classes of destinations - local housing, MARTA, or a parking facility. While an important simplification, it is likely true in reality that pedestrians may be traveling to other destinations as well, such as social venues. It is also likely that some percentage of pedestrians may not evacuate the area for an extended time period, as they sit idle while waiting for friends or planning their destination in an ad-hoc manner.

We assume the stochastic processes in this simulation to be homogenous stochastic processes, which can be defined as stationary stochastic processes where the random variables $X = X_1, X_2, \dots, X_n$ are independent and identically distributed. While not always realistic in the context of simulation model-building, we do believe this assumption to be reasonable in the context of this simulation.

There are inherent simplifications in the choice of a cellular automata-based simulation approach. Specifically, it is assumed that each individual travels a constant distance per time step through the simulation. Nevertheless, we attempt to mitigate this simplification by drawing the speed for each individual from a probability distribution, which does introduce some variation, albeit only on an individual-to-individual basis.

4 Description of Simulation Software

4.1 Architecture

In generating our simulation, we take an object-oriented design approach. We use the Python programming language for our implementation.

Distinct types of entities in our simulation are represented as *classes*. For interacting with these classes, we define a number of methods that function as their public interfaces. It should be noted that while we may define explicit “getter” and “setter” methods in our conceptual model, the Python programming language provides these interfaces on declared object attributes without the need to provide explicit method definitions in code.

4.2 Interfaces

4.2.1 User Interface

To run our simulation, first ensure Python 2.7 is installed on the testing machine. Python’s matplotlib module will also need to be installed for visualization purposes. If visualization is not needed, it can be turned off by setting the *Visualization* option to False when initializing the simulation.

Python’s NumPy module will also be necessary for random number sampling from the Poisson distribution. In addition, a configuration JSON file must be provided to configure the intersections. A sample configuration JSON file is shown below, and all the configuration files used for testing our simulation model are provided in the *code/config* directory.

```
{  
    "name": "config2",  
    "num_sims": "20",  
    "parameters": [  
        {  
            "name": "Parameter 1 - closed intersections",  
            "type": "intersection_closed",  
            "data": {  
                "intersections": [  
                    "3",  
                    "9"  
                ]  
            }  
        }  
    ]  
}
```

```

},
{
  "name": "Parameter 2 - open intersections",
  "type": "intersection_open",
  "data": {
    "intersections": [
      "2"
    ]
  }
},
{
  "name": "Parameter 3 - normal intersections",
  "type": "intersection_normal",
  "data": {
    "intersections": [
      {
        "id": "1",
        "time": "15"
      },
      {
        "id": "4",
        "time": "30"
      },
      {
        "id": "5",
        "time": "25"
      },
      {
        "id": "6",
        "time": "20"
      },
      {
        "id": "7",
        "time": "15"
      },
      {
        "id": "8",
        "time": "15"
      }
    ]
  }
}

```

```

        "id": "10",
        "time": "15"
    },
    {
        "id": "11",
        "time": "40"
    },
    {
        "id": "12",
        "time": "25"
    }
]
}
]
}

```

Within the configuration above, there is a *name* field which specifies a label for the configuration, *num_sims* field which specifies the number of simulations that should be run, and *parameters* array, which contains three types of parameters - *intersection_closed*, *intersection_open*, and *intersection_normal*. For each of these, a list of intersection identifiers is given. These identifiers correspond to those in Fig. 5, and instruct the simulation on whether to close the intersections to pedestrians, open the intersections to pedestrians, or leave the intersection as normal - in other words, as changing between closed and open at some interval.

For the *intersection_normal* case, instead of a simple list of intersection identifiers, a list of objects should be given as shown above, with a *time* parameter specifying the number of timesteps each intersection should be closed.

Once Python, matplotlib, and NumPy are installed and a configuration JSON file is selected, you can simply enter the *code* subdirectory and run the following at a Unix command prompt to start an example batch of 20 simulation runs with 5,000 pedestrians:

```
$ python sim_batch.py -c config/config1.json -p 5000 -v t -f
  paths/config1.pickle
```

The more general form is as follows:

```
$ python sim_batch.py -c <configJsonFile> -p <int numPeds> -v <t/f  
vizBoolean> -f <pathsFile>
```

where *configJsonFile* is the location of your configuration JSON file, *numPeds* is an integer defining the number of pedestrians to simulate, *vizBoolean* is *t* or *f* stating whether or not visualization should be turned on or off, respectively, and *pathsFile* is the location of a *pickle* file, which contains precomputed shortest path information.

The output will initially state that a preprocessing step is being performed to prepare the simulation. Then, the actual simulation will automatically begin. Pedestrians will be created and will begin moving toward their destinations. As they do, the number of “active pedestrians” remaining in the SUI and the number of pedestrians remaining in the pedestrian queue (in other words, the input queue) will be displayed at every time step.

Behind the scenes, initializing a new simulation creates a *Grid* object, which is initialized from a node file, intersection file, edge file, type map giving translations between integer node types (such as SUI entrance, SUI exit, sidewalk, and road), and configuration file containing parameterizations. An optional file in *pickle* format can contain previously computed shortest path information, as this is a fairly expensive step that need only occur when intersections are opened or closed.

The node file should be in a CSV format, of the form:

```
cell_x, cell_y, pixel_x, pixel_y, meters_x, meters_y, node_type
```

where *cell_x* and *cell_y* correspond to *x* and *y* coordinates in a 2D-grid space. *pixel_x*, *pixel_y*, *meters_x*, and *meters_y* translate the *x* and *y* coordinates into pixel space and physical space, respectively. The *node_type* is an integer value that specifies whether the node is a sidewalk (1), road (2), SUI entrance (3), or SUI exit (4) node.

The intersection file should also be in a CSV format, of the form:

```
intersection_id, node_id
```

where the given node id represents a node that is part of the intersection with the corresponding id.

The edge file should also be in a CSV format, of the form:

```
node_a_id, node_b_id, edge_weight
```

where the node ids refer to zero-based line indices in the node file. In other words, if the node file contained two rows, such as:

```
1,4,4,16,4,16,2  
2,6,8,24,8,24,2
```

A corresponding edge file might look as follows:

```
0,1,7.000
```

where 0 refers to the node in the first line of the nodes file, and 1 refers to the node in the second line. The type map should be a dictionary, with string keys that are human-readable descriptions of node types, and integer values that correspond to node types in the node file.

Once the grid is created, it is passed to the *Simulation* constructor, along with a parameter dictionary with options such as the number of pedestrians that should be simulated, and whether the visualization engine will be enabled. An overview of the key relationships between classes in our simulation is shown in Fig. 6.

4.2.2 Pedestrian Class Interfaces

For simulating pedestrian movement, we define a *Pedestrian* class as shown in Table 5. Pedestrians are initialized with an entrance node to the SUI, as well as a destination node and speed. They then traverse the nodes in the shortest path toward their destination node. Once they have reached a node with an *exit* NodeType, the SetEgressComplete() method is called for the pedestrian with an argument of true, alerting the simulation that the pedestrian has finished the simulation. Once there are no pedestrians in the simulation remaining with the EgressComplete attribute set to false, the driver program breaks from the main simulation loop and the program exits.

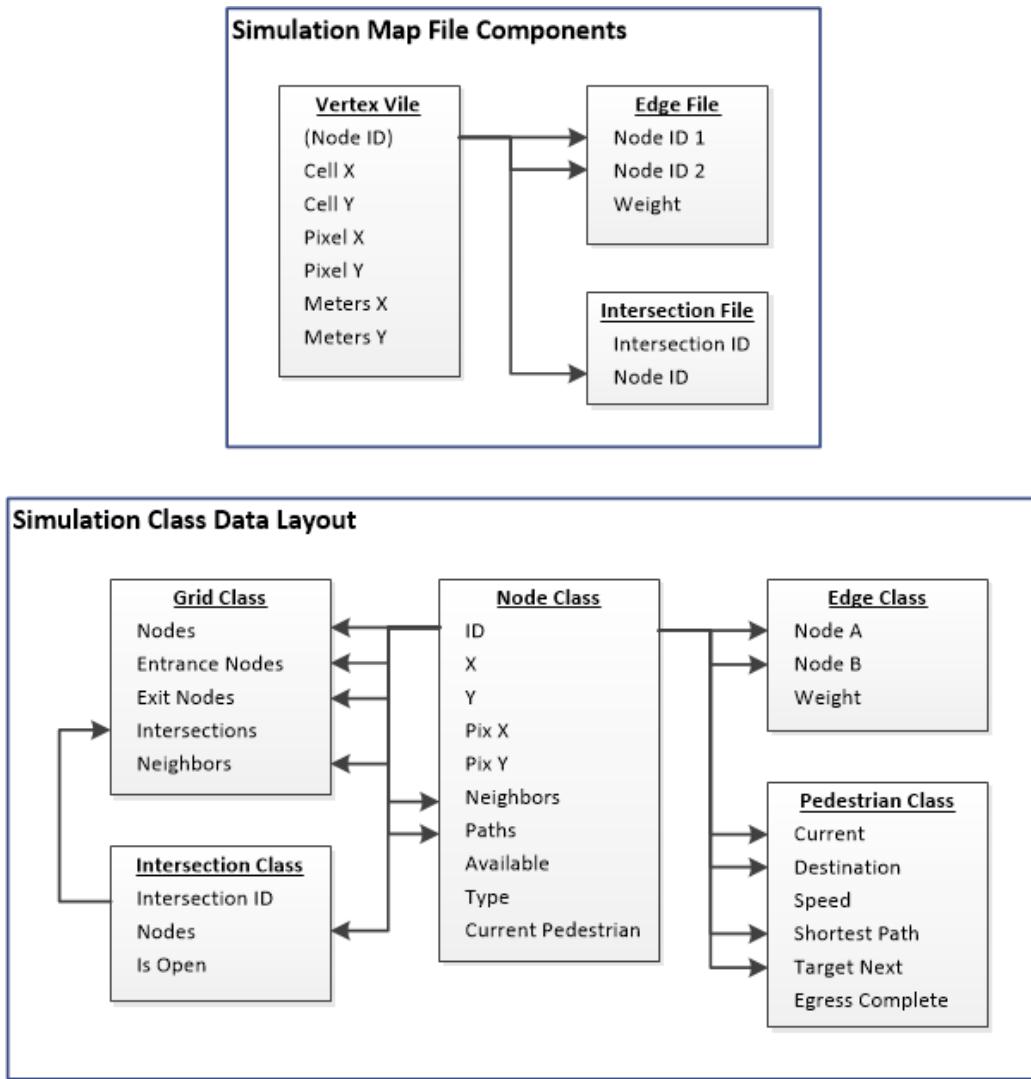


Figure 6: Simulation Data Layout.

4.2.3 Node Class Interfaces

We define a *Node* class as shown in Table 6. Nodes are the primary abstraction for the resource entities in the simulation, which are sidewalks and roads. Nodes have several key attributes defined that can be retrieved and set. Among them, the *Neighbors* attribute specifies all the nodes that are

Method	Description
Move(NewNode)	Takes one argument: a <i>Node</i> object. When called, marks the node set to <i>Current</i> as unavailable, sets the <i>Current</i> attribute to NewNode, and marks the NewNode object as unavailable.
GetCurrentNode() SetCurrentNode(CurrentNode)	Returns or sets the current node that the pedestrian is occupying.
GetDestinationNode() SetDestinationNode(DestNode)	Returns or sets the destination node for the pedestrian.
GetTargetNext() SetTargetNext(TargetNext)	Returns or sets the target next node for the pedestrian in the shortest path to his destination.
GetEgressComplete() SetEgressComplete(Complete)	Returns or sets the current egress status of the pedestrian. True once the pedestrian has exited the SUI, and false otherwise.

Table 5: Methods of the *Pedestrian* class.

Method	Description
GetAvailable() SetAvailable(Bool)	Gets or sets the current value of the <i>Available</i> attribute for the node. If the setter, must be a boolean true or false.
GetPaths() SetPaths(Paths)	Returns or sets the <i>Paths</i> attribute of the node, which is a hash map that gives the next node in the shortest path to each possible destination node in the simulation.
GetNeighborNodes() SetNeighborNodes(ListNodes)	Returns or sets the <i>Neighbors</i> attribute of the node, corresponding to the nodes one edge length away. If the setter, takes one argument: list of <i>Node</i> objects.
GetNodeType()	Returns the <i>NodeType</i> Attribute of the node.
GetNodeID()	Returns the <i>NodeID</i> Attribute of the node.

Table 6: Methods of the *Node* class.

one degree away from the node object. The *Neighbors* attribute is then used to populate the *Paths* attribute for the node, which is a hash map that gives the shortest path to each possible destination node in the simulation. We precompute the *Paths* attribute for each entrance node in the graph prior to beginning the simulation. This precomputation step needs only be repeated when a parameter change occurs that alters the underlying node structure or availability.

4.2.4 Intersection Class Interfaces

Within our simulation, we utilize a number of intersections that are parametrically closed and opened. Thus, we define key methods for the *Intersection* class in Table 7.

Intersections are used to limit the rate of flow of pedestrians through the

Method	Description
OpenMe()	When called, sets the <i>Available</i> attribute to true for every node in its <i>Nodes</i> array. Sets the <i>Open</i> attribute for the intersection to true.
CloseMe()	When called, sets the <i>Available</i> attribute to false for every node in its <i>Nodes</i> array. Sets the <i>Open</i> attribute for the intersection to false.
GetNodes()	Returns the list of <i>Node</i> objects that are part of the intersection.
SetNodes(ListNodes)	Takes one argument: list of <i>Node</i> objects. Sets the <i>Node</i> objects that make up the intersection.

Table 7: Methods of the *Intersection* class.

SUI, increasing the reality of the simulation. As an intersection contains one or more nodes, when an intersection is closed, all the nodes within it are closed to additional pedestrian traffic. This means no new pedestrian traffic is allowed into the intersection. Pedestrians currently in the intersection as it closes will not be able to move if their next node is within the intersection. Although keeping pedestrians trapped in the middle of the road is unrealistic, it simplifies the model and the effect on the simulation was decided to be negligible as the same “delay” effect is still executed.

5 Results

As mentioned earlier, several intersections were identified in the SUI. These intersections are shown in Fig. 5.

Ten distinct combinations of intersection states were created with the purpose of evaluating the time it takes for pedestrians to evacuate the SUI under each combination; this is shown in Table 8. Green cells indicate an intersection open exclusively for pedestrians and closed to vehicular traffic (i.e., it is effectively a sidewalk), yellow cells indicate an intersection with

	Configuration #									
	1	2	3	4	5	6	7	8	9	10
Intersection ID	1	open	normal	open	normal	open	normal	open	normal	open
2	normal	open	open	normal	normal	open	open	normal	normal	normal
3	normal	closed	normal	normal	normal	closed	open	normal	closed	closed
4	normal	normal	normal	normal	normal	normal	normal	normal	normal	normal
5	open	normal	normal	normal	open	normal	open	open	closed	open
6	open	normal	closed	normal	open	normal	open	open	normal	open
7	normal	normal	normal	normal	normal	normal	normal	normal	normal	normal
8	normal	normal	normal	normal	normal	normal	normal	normal	normal	normal
9	normal	closed	normal	normal	closed	normal	normal	normal	closed	normal
10	normal	normal	normal	normal	normal	normal	normal	normal	normal	normal
11	normal	normal	normal	normal	normal	normal	normal	normal	normal	normal
12	normal	normal	normal	normal	normal	normal	normal	normal	normal	normal

Table 8: Intersection Configurations

alternating vehicular and pedestrian traffic (i.e., a law enforcement officer directing traffic or a normal crosswalk), and red cells indicate an intersection closed to all pedestrian traffic (i.e., nodes are removed from the map).

Intersection configurations were carefully configured to adhere to two governing rules:

1. Intersections on North Ave shall never be open (i.e., closed to vehicular traffic).
2. No closed intersections shall make any destination node unreachable.

To achieve statistical significance, 20 runs were performed per intersection configuration. As stated earlier, 90% confidence intervals were constructed from each set of 20 runs. The plot of the expected value, along with error bars representing the 90% confidence intervals, is shown in Fig. 7. The runs are sorted such that the lowest expected evacuation times increase from left to right.

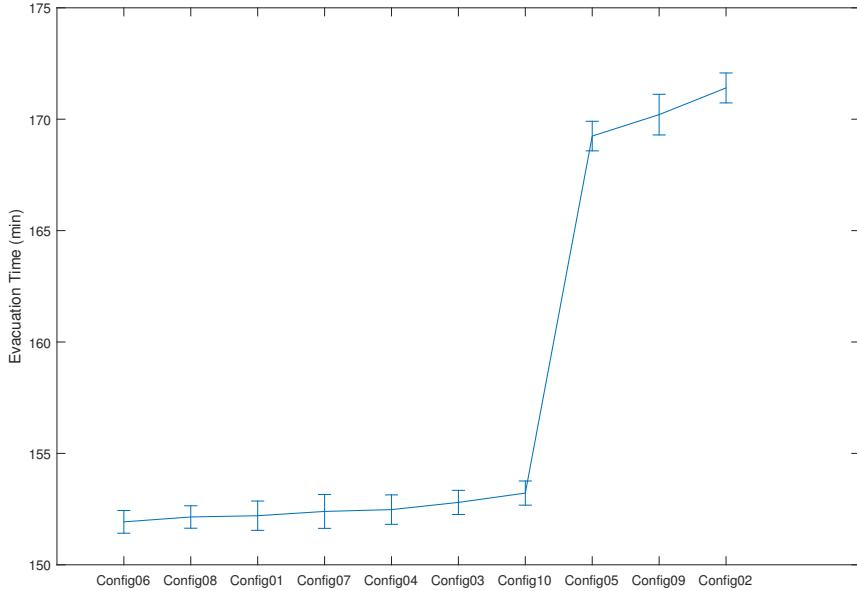


Figure 7: SUI Evacuation Times.

There are a few takeaways from the parametric results above. First, the step-like behavior in evacuation times is caused exclusively by the closure of intersection 9; this is one of the few pathways that cross North Avenue. According to the results, its closure leads to an increase in average evacuation time by roughly 15 minutes. The second takeaway is that, aside from the state of intersection 9 (which is never green, as this would violate the rule of not closing North Avenue to vehicular traffic), the state of other intersections do not have a statistically significant impact to evacuation time. Finally, the 90% confidence intervals are not very large: the maximum interval was merely 0.9118 minutes, indicating the expected values plotted are, at worst, 90% likely to be within about 1 minute of the true mean of each configuration's random variable.

An example of the impact of closing intersection 9 follows. Consider configuration 2 and configuration 6; note that the only difference between these two configurations is the state of intersection 9. Similarly, configuration 5 and configuration 1 exhibit this same relationship. In Fig. 7, note that configurations 2 and 5 turned out to be on the upper portion of the step-like

function, while configurations 6 and 1 are on the lower portion; this lends further validity to the notion that the closure of intersection 9 is indeed the only culprit (in the set of tested intersection configurations) of longer evacuation times observed in the simulation runs.

5.1 Verification Methodology

Throughout development of the model, team members routinely reviewed code that they did not write; this ensured that outside perspectives and fresh sets of eyes reviewed the code. Code is very modular and compartmented to avoid a multitude of programmatic pitfalls.

Before running on the SUI depicted in Fig. 2, a smaller map was generated to evaluate model execution on a much smaller scale (the smaller map featured 2,380 nodes, while the SUI investigated in this study contained 107,498 nodes). This facilitated rapid identification and fixing of bugs in a scalable simulation, and instilled confidence that the simulation was running as expected on the larger map.

Visualization was a critical aspect of verifying the simulation model, and the development team regularly reviewed recordings of visualizations, in addition to carrying out real-time visual inspection of runs as they happened.

Addressing the soundness of the vertex, edge, and intersection files, sanity checks were performed to avoid common pitfalls (such as edges referencing nodes that do not exist and intersections referencing nodes not representing roadways). One-wayness of map file generation also precluded tool-related bugs if one wanted to change these entities.

As described earlier, goodness-of-fit tests were employed to ensure random number generation was satisfactorily performed.

5.2 Validation Methodology

Literature of comparable pedestrian simulations with similar outputs (i.e., evacuation time analyses) was extensively used to define details regarding our model. A comprehensive description of takeaways from literature may be found in Section 2.

As specified earlier, the pedestrian interarrival time was derived experimentally by observing a crowd of pedestrians exiting a structure. The data from this was scaled to meet the size of the exits in our particular SUI.

The evacuation times themselves were evaluated based upon each author's experiences, which provides a level of face validation. In addition, configurations expected to have higher evacuation times (such as configuration 9, which had the most closed intersections) and lower evacuation times (configuration 7, with the most open intersections) were identified, and the results aligned with this intuition.

Again, visualization was utilized to ensure the simulation of pedestrian motion matched the authors' expectations based on real-world experiences.

Finally, only practical intersection configurations were evaluated, where a practical intersection is defined as adhering to the governing rules listed earlier.

6 Conclusion

While the closure of intersection 9 (shown in Fig. 5) to pedestrians appears to be responsible for an approximate 15 minute increase in evacuation time (raising the evacuation time from 2.5 hours to 2.75 hours), this amounts to an approximate 10% increase in evacuation time from the SUI as compared to when intersection 9 is operating like a conventional crosswalk.

For the intersection configurations tested, there was not a single configuration that stood out as a configuration that yields significantly lower evacuation times relative to the other configurations. As such, so long as pedestrians are able to arrive at their destinations, under the assumptions noted in this paper, there is no evidence of a superior configuration which dramatically reduces evacuation time.

7 Final Notes

It should be noted that all of the authors contributed equally to this work. From a testing perspective, our simulation model has been tested primarily in a Mac OS X environment - specifically, OS X 10.11.3 with Python 2.7.

Thanks are due to the TAs and instructor for their assistance with our many questions.

References

- [1] Ofer Biham, A Alan Middleton, and Dov Levine. Self-organization and a dynamical transition in traffic-flow models. *Physical Review A*, 46(10):R6124, 1992.
- [2] Victor J Blue and Jeffrey L Adler. Cellular automata microsimulation for modeling bi-directional pedestrian walkways. *Transportation Research Part B: Methodological*, 35(3):293–312, 2001.
- [3] Carsten Burstedde, Kai Klauck, Andreas Schadschneider, and Johannes Zittartz. Simulation of pedestrian dynamics using a two-dimensional cellular automaton. *Physica A: Statistical Mechanics and its Applications*, 295(3):507–525, 2001.
- [4] Edsger W Dijkstra. A note on two problems in connexion with graphs. *Numerische mathematik*, 1(1):269–271, 1959.
- [5] David Eppstein. Dijkstra’s algorithm for shortest paths. <https://www.ics.uci.edu/~eppstein/161/python/dijkstra.py>, 2002.
- [6] Zhixiang Fang, Xinlu Zong, Qingquan Li, Qiuping Li, and Shengwu Xiong. Hierarchical multi-objective evacuation routing in stadium using ant colony optimization approach. *Journal of Transport Geography*, 19(3):443–451, 2011.
- [7] Dirk Helbing. A fluid dynamic model for the movement of pedestrians. *arXiv preprint cond-mat/9805213*, 1998.
- [8] Dirk Helbing, Lubos Buzna, Anders Johansson, and Torsten Werner. Self-organized pedestrian crowd dynamics: Experiments, simulations, and design solutions. *Transportation science*, 39(1):1–24, 2005.
- [9] Dirk Helbing, Illés Farkas, and Tamas Vicsek. Simulating dynamical features of escape panic. *Nature*, 407(6803):487–490, 2000.
- [10] Armel Ulrich Kemloh Wagoum, Armin Seyfried, and Stefan Holl. Modeling the dynamic route choice of pedestrians to assess the criticality of building evacuation. *Advances in Complex Systems*, 15(07):1250029, 2012.

- [11] Hubert Klüpfel, Michael Schreckenberg, and Tim Meyer-König. Models for crowd movement and egress simulation. In *Traffic and Granular Flow'03*, pages 357–372. Springer, 2005.
- [12] Shigeyuki Okazaki and Satoshi Matsushita. A study of simulation model for pedestrian movement with evacuation and queuing. In *International Conference on Engineering for Crowd Safety*, pages 271–280, 1993.
- [13] Stephen K. Park and Keith W. Miller. Random number generators: good ones are hard to find. *Communications of the ACM*, 31(10):1192–1201, 1988.
- [14] WH Payne, John R Rabung, and TP Bogyo. Coding the lehmer pseudo-random number generator. *Communications of the ACM*, 12(2):85–86, 1969.
- [15] PhysicsWorld. Pedestrian-dynamics experiment: stadium evacuation. <https://www.youtube.com/watch?v=4AZQ4lFLcb4>, 2010.
- [16] Khalidur Rahman, Noraida Abdul Ghani, Anton Abdulbasah Kamil, Adli Mustafa, and Md Ahmed Kabir Chowdhury. Modelling pedestrian travel time and the design of facilities: A queuing approach. *PloS one*, 8(5):e63503, 2013.
- [17] Fang Weifeng, Yang Lizhong, and Fan Weicheng. Simulation of bi-direction pedestrian movement using a cellular automata model. *Physica A: Statistical Mechanics and its Applications*, 321(3):633–640, 2003.
- [18] Douglas Brent West et al. *Introduction to graph theory*, volume 2. Prentice hall Upper Saddle River, 2001.
- [19] Xiaoping Zheng, Tingkuan Zhong, and Mengting Liu. Modeling crowd evacuation of a building based on seven methodological approaches. *Building and Environment*, 44(3):437–445, 2009.